

house-price-prediction

January 17, 2026

IMPORTING NECESSARY LIBS

```
[517]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab import drive
drive.mount('/content/drive') # Mount Google Drive
df=pd.read_csv('/content/drive/MyDrive/House Price Prediction Dataset.csv')
df.head()
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call
drive.mount("/content/drive", force_remount=True).

```
[517]:   Id Area Bedrooms Bathrooms Floors YearBuilt Location Condition \
0    1  1360          5         4      3    1970 Downtown Excellent
1    2  4272          5         4      3    1958 Downtown Excellent
2    3  3592          2         2      3    1938 Downtown     Good
3    4    966          4         2      2    1902 Suburban    Fair
4    5  4926          1         4      2    1975 Downtown    Fair

   Garage    Price
0      No  149919
1      No  424998
2      No  266746
3     Yes  244020
4     Yes  636056
```

DATA CLEANING AND PREPROCESSING

```
[518]: df.shape
#no.of rows=2000 and columns = 10
```

[518]: (2000, 10)

```
[519]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
```

```
Data columns (total 10 columns):
 #   Column      Non-Null Count Dtype  
 --- 
 0   Id          2000 non-null    int64  
 1   Area         2000 non-null    int64  
 2   Bedrooms     2000 non-null    int64  
 3   Bathrooms    2000 non-null    int64  
 4   Floors        2000 non-null    int64  
 5   YearBuilt    2000 non-null    int64  
 6   Location      2000 non-null    object  
 7   Condition     2000 non-null    object  
 8   Garage        2000 non-null    object  
 9   Price         2000 non-null    int64  
dtypes: int64(7), object(3)
memory usage: 156.4+ KB
```

```
[520]: df.isnull().sum()
#No null value found
```

```
[520]: Id          0
Area         0
Bedrooms     0
Bathrooms    0
Floors        0
YearBuilt    0
Location      0
Condition     0
Garage        0
Price         0
dtype: int64
```

```
[521]: df.describe()
```

```
[521]:           Id          Area      Bedrooms      Bathrooms      Floors \ 
count  2000.000000  2000.000000  2000.000000  2000.000000  2000.000000
mean   1000.500000  2786.209500  3.003500    2.552500   1.993500
std    577.494589  1295.146799  1.424606    1.10899   0.809188
min    1.000000    501.000000   1.000000    1.00000   1.000000
25%   500.750000  1653.000000  2.000000    2.00000   1.000000
50%   1000.500000  2833.000000  3.000000    3.00000   2.000000
75%   1500.250000  3887.500000  4.000000    4.00000   3.000000
max   2000.000000  4999.000000  5.000000    4.00000   3.000000

           YearBuilt      Price
count  2000.000000  2000.000000
mean   1961.446000  537676.855000
std    35.926695   276428.845719
```

```
min      1900.000000  50005.000000
25%     1930.000000  300098.000000
50%     1961.000000  539254.000000
75%     1993.000000  780086.000000
max     2023.000000  999656.000000
```

```
[522]: df.duplicated().sum()
#No duplicate value found
```

```
[522]: np.int64(0)
```

```
[523]: df.columns
```

```
[523]: Index(['Id', 'Area', 'Bedrooms', 'Bathrooms', 'Floors', 'YearBuilt',
       'Location', 'Condition', 'Garage', 'Price'],
       dtype='object')
```

MOVING UNNECESSARY COLUMNS

```
[524]: df=df.drop(columns=[ 'Bedrooms', 'Bathrooms','Id'])
df.head()
```

```
[524]:   Area  Floors  YearBuilt  Location  Condition  Garage  Price
0    1360      3        1970  Downtown  Excellent    No  149919
1    4272      3        1958  Downtown  Excellent    No  424998
2    3592      3        1938  Downtown      Good    No  266746
3     966      2        1902 Suburban      Fair   Yes  244020
4    4926      2        1975  Downtown      Fair   Yes  636056
```

```
[525]: df['Area'].unique()
```

```
[525]: array([1360, 4272, 3592, ..., 865, 2174, 4062])
```

```
[526]: df['Location'].unique()
```

```
[526]: array(['Downtown', 'Suburban', 'Urban', 'Rural'], dtype=object)
```

```
[527]: df['Location'].isnull().sum()
```

```
[527]: np.int64(0)
```

```
[528]: df['Floors'].unique()
```

```
[528]: array([3, 2, 1])
```

```
[529]: df['Garage'].unique()
```

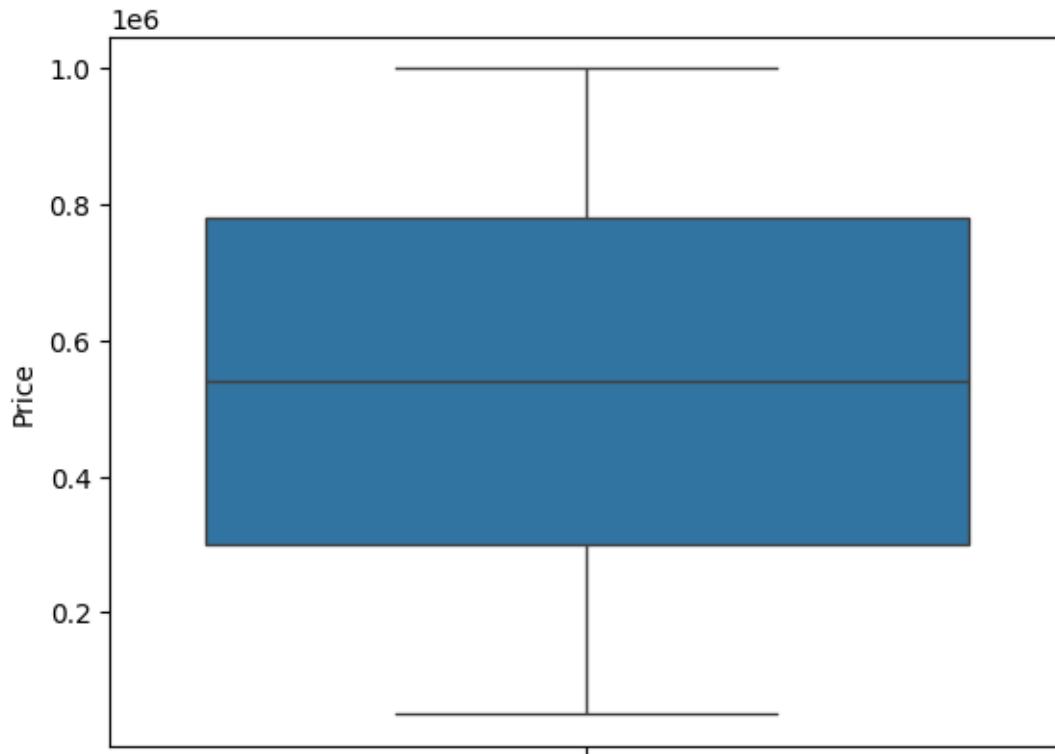
```
[529]: array(['No', 'Yes'], dtype=object)
```

```
[530]: df['YearBuilt'].unique()
```

```
[530]: array([1970, 1958, 1938, 1902, 1975, 1906, 1948, 1925, 1932, 2000, 1947,  
1978, 1901, 2004, 1931, 1903, 1919, 2013, 2016, 1935, 1927, 1976,  
1900, 1959, 1955, 1934, 2011, 1929, 1953, 2020, 1954, 1988, 1979,  
1957, 1982, 1964, 1968, 1950, 1921, 1987, 2006, 2008, 2015, 1952,  
1999, 1967, 1951, 1981, 1949, 1940, 1917, 1965, 1920, 1943, 2002,  
1946, 1928, 1989, 1984, 1916, 1930, 2014, 1972, 1994, 1977, 2009,  
1913, 1996, 1998, 2010, 1983, 2022, 1915, 1911, 2018, 1904, 1980,  
2021, 2005, 1973, 1942, 1944, 1908, 1961, 1956, 1924, 1914, 1905,  
2019, 1941, 1992, 1974, 1963, 2001, 1991, 1936, 1907, 1997, 2007,  
2017, 1966, 1945, 1912, 1986, 1960, 1995, 1933, 1969, 1923, 2012,  
1910, 2003, 1993, 2023, 1918, 1971, 1926, 1939, 1922, 1937, 1909,  
1990, 1962, 1985])
```

```
[531]: sns.boxplot(df['Price'])
```

```
[531]: <Axes: ylabel='Price'>
```



```
[532]: df['Price'].max()
```

```
[532]: 999656
```

```
[533]: df['Price'].min()
```

```
[533]: 50005
```

```
[534]: df['Price'].mean()
```

```
[534]: np.float64(537676.855)
```

DATA EXTRACTION

```
[535]: x=df[['Area', 'Floors', 'YearBuilt', 'Location', 'Condition', 'Garage']]  
y=df['Price']  
x
```

```
[535]:      Area  Floors  YearBuilt  Location  Condition  Garage  
0       1360      3       1970  Downtown  Excellent    No  
1       4272      3       1958  Downtown  Excellent    No  
2       3592      3       1938  Downtown     Good     No  
3        966      2       1902 Suburban     Fair     Yes  
4       4926      2       1975  Downtown     Fair     Yes  
...     ...     ...     ...     ...     ...  
1995    4994      3       1923 Suburban     Poor     No  
1996    3046      1       2019 Suburban     Poor     Yes  
1997    1062      2       1903   Rural     Poor     No  
1998    4062      2       1936   Urban  Excellent     Yes  
1999    2989      3       1903 Suburban     Fair     No
```

[2000 rows x 6 columns]

```
[536]: y
```

```
[536]: 0      149919  
1      424998  
2      266746  
3      244020  
4      636056  
...  
1995    295620  
1996    580929  
1997    476925  
1998    161119  
1999    482525  
Name: Price, Length: 2000, dtype: int64
```

APPLYING TRAIN SPLIT

```
[537]: #IMPORTING LIBRARY  
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)  
x_train
```

```
[537]:      Area  Floors  YearBuilt  Location  Condition  Garage  
368    4230       3        2018  Suburban       Good     No  
488    2714       1        1966   Urban       Good    Yes  
657    3378       3        1998  Rural  Excellent     No  
1699   3041       2        1987 Downtown       Fair     No  
338    1581       3        2006  Rural  Excellent     No  
...     ...     ...     ...     ...     ...  
810    4518       1        1998   Urban  Excellent     No  
1560   4068       2        1956   Urban  Excellent     No  
1843   3767       1        1901 Downtown       Fair     No  
686    3783       2        1961 Suburban       Good    Yes  
887    2354       3        2007 Suburban  Excellent     No
```

[1400 rows x 6 columns]

```
[538]: x_test
```

```
[538]:      Area  Floors  YearBuilt  Location  Condition  Garage  
1180   2734       3        2009  Suburban  Excellent     No  
1271   1325       2        2008  Rural       Poor     No  
1631   4328       2        1970 Suburban  Excellent    Yes  
883    1705       1        1985  Rural       Fair    Yes  
946    4650       2        1923  Rural       Fair    Yes  
...     ...     ...     ...     ...     ...  
1711   1786       1        1945   Urban       Fair     No  
1973   4526       3        1936  Rural       Fair     No  
728    503        3        2012 Suburban       Poor     No  
613    1858       1        2013  Rural       Good     No  
1007   1804       1        1991   Urban       Good    Yes
```

[600 rows x 6 columns]

ENCODING NOMINAL DATA

```
[539]: from sklearn.preprocessing import OneHotEncoder,StandardScaler  
ohe=OneHotEncoder()  
ohe.fit(x[['Location','Garage']])
```

```
[539]: OneHotEncoder()
```

ORDINAL ENCODING

```
[540]: from sklearn.preprocessing import OrdinalEncoder  
oe=OrdinalEncoder()  
oe.fit(x[['Condition']])
```

```
[540]: OrdinalEncoder()
```

SCALING DATA

```
[541]: scaler=StandardScaler()  
scaler.fit(x[['Area']])
```

```
[541]: StandardScaler()
```

```
[542]: df['Condition'].unique()
```

```
[542]: array(['Excellent', 'Good', 'Fair', 'Poor'], dtype=object)
```

COLUMN TRANSFORMATION

```
[543]: from sklearn.compose import make_column_transformer  
column_transform=make_column_transformer((OneHotEncoder(categories=ohe.  
↳categories_),['Location','Garage']),  
↳(OrdinalEncoder(categories=[[('Excellent', 'Good', 'Fair',  
↳'Poor')]],['Condition']),  
↳(StandardScaler(),['Area']),  
remainder='passthrough')  
column_transform
```

```
[543]: ColumnTransformer(remainder='passthrough',  
transformers=[('onehotencoder',  
OneHotEncoder(categories=[array(['Downtown',  
'Rural', 'Suburban', 'Urban'], dtype=object),  
array(['No', 'Yes'],  
dtype=object)],  
['Location', 'Garage']),  
('ordinalencoder',  
OrdinalEncoder(categories=[[('Excellent',  
'Good', 'Fair',  
'Poor')]],  
['Condition']),  
('standardscaler', StandardScaler(), ['Area']))])
```

```
[544]: x_train_trans=column_transform.fit_transform(x_train)  
x_test_trans=column_transform.transform(x_test)
```

MODELING

```
[545]: #LINEAR REGRESSION  
from sklearn.linear_model import LinearRegression  
  
lr=LinearRegression()
```

```

#lr.fit(x_train_trans,y_train)

#TREE BASED REGRESSION
from sklearn.tree import DecisionTreeRegressor
dt=DecisionTreeRegressor()
#dt.fit(x_train_trans,y_train)

#DISTANCE BASED MODELS
from sklearn.neighbors import KNeighborsRegressor
knn=KNeighborsRegressor()
#knn.fit(x_train_trans,y_train)

#BOOSTING
from lightgbm import LGBMRegressor
lgb=LGBMRegressor()
lgb

```

[545]: LGBMRegressor()

Linear Regression: It imports the LinearRegression model from sklearn.linear_model, creates an instance named lr, and then trains (fits) this model using your transformed training features (x_train_trans) and the corresponding training prices (y_train). Decision Tree Regressor: Similarly, it imports DecisionTreeRegressor from sklearn.tree, creates an instance dt, and fits it to the same training data. Decision trees are non-linear models that learn decision rules from data. K-Neighbors Regressor: This part imports KNeighborsRegressor from sklearn.neighbors, creates an instance knn, and fits it. This is a distance-based model that predicts the target based on the ‘k’ nearest data points in the training set. LGBM Regressor (Light Gradient Boosting Machine): Finally, it imports LGBMRegressor from lightgbm and creates an instance lgb. LightGBM is a gradient boosting framework that uses tree-based learning algorithms. However, this specific line only initializes the lgb object; it does not train (fit) the lgb model in this cell.

[546]: regressors={
 'LR':lr,
 'DT':dt,
 'KNN':knn,
 'LGB':lgb
 }

This code block creates a Python dictionary called regressors. It’s essentially a container to store the different machine learning models you’ve initialized earlier, making them easily accessible using simple string keys.

‘LR’ (Linear Regression) is mapped to the lr model instance. ‘DT’ (Decision Tree Regressor) is mapped to the dt model instance. ‘KNN’ (K-Neighbors Regressor) is mapped to the knn model instance. ‘LGB’ (LightGBM Regressor) is mapped to the lgb model instance. This dictionary structure is very useful for iterating through your models, making predictions with each, and evaluating their performance.

TRAINING THE REGRESSOR

```
[547]: def train_regressor(regressor,x_train,y_train,x_test,y_test,name):
    regressor.fit(x_train,y_train)
    regressor
    y_pred=regressor.predict(x_test)
    y_pred
```

CALCULATING METRICS

```
[548]: #IMPORTING LIB FOR METRICS
from sklearn.metrics import r2_score,mean_absolute_error,mean_squared_error

metrics_results = {}

for name, regressor in regressors.items():
    print(f"Evaluating {name}:")
    # Fit the regressor before making predictions
    regressor.fit(x_train_trans, y_train)

    # Make predictions on the transformed test set
    y_pred = regressor.predict(x_test_trans)

    # Calculate metrics
    mse = mean_squared_error(y_test, y_pred)
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    metrics_results[name] = {
        'MSE': mse,
        'MAE': mae,
        'R2': r2
    }
    print(f"  MSE: {mse:.2f}")
    print(f"  MAE: {mae:.2f}")
    print(f"  R2 Score: {r2:.2f}")
    print("\n")

# Optional: Display all results at once
# import pandas as pd
# print(pd.DataFrame(metrics_results).T)
```

Evaluating LR:

```
MSE: 73973479746.08
MAE: 237431.98
R2 Score: 0.00
```

Evaluating DT:

```
MSE: 154873056519.03
```

```
MAE: 325111.43
R2 Score: -1.09
```

```
Evaluating KNN:
MSE: 91532710459.61
MAE: 255885.91
R2 Score: -0.24
```

```
Evaluating LGB:
[LightGBM] [Info] Auto-choosing col-wise multi-threading, the overhead of
testing was 0.000181 seconds.
You can set `force_col_wise=true` to remove the overhead.
[LightGBM] [Info] Total Bins 400
[LightGBM] [Info] Number of data points in the train set: 1400, number of used
features: 10
[LightGBM] [Info] Start training from score 533118.275000
MSE: 86904992222.44
MAE: 251961.90
R2 Score: -0.17
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMRegressor was fitted
with feature names
    warnings.warn(
```

This code block is crucial for evaluating the performance of your trained regression models. Here's what it does step-by-step:

Import Metrics Libraries: It starts by importing `r2_score`, `mean_absolute_error`, and `mean_squared_error` from `sklearn.metrics`. These functions are used to quantify how well your models perform.

Initialize metrics_results Dictionary: An empty dictionary named `metrics_results` is created. This dictionary will store the calculated evaluation metrics for each of your models.

Iterate Through Models: The code then loops through each model stored in the `regressors` dictionary (which you defined earlier, containing `lr`, `dt`, `knn`, and `lgb`). In each iteration, `name` gets the model's key (e.g., '`LR`', '`DT`') and `regressor` gets the actual trained model object.

Make Predictions: Inside the loop, `y_pred = regressor.predict(x_test_trans)` uses the current trained regressor to make predictions on your transformed test features (`x_test_trans`). `y_pred` will be an array of predicted house prices.

Calculate Metrics:

`mse = mean_squared_error(y_test, y_pred)`: Calculates the Mean Squared Error, which measures the average of the squares of the errors. Larger errors are penalized more heavily. `mae = mean_absolute_error(y_test, y_pred)`: Calculates the Mean Absolute Error, which is the average

of the absolute differences between the actual and predicted values. It gives a more direct sense of the average prediction error. `r2 = r2_score(y_test, y_pred)`: Calculates the R-squared (coefficient of determination) score. This indicates the proportion of the variance in the target variable that your model can explain. A higher R2 score (closer to 1) means a better fit. Store and Print Results: The calculated mse, mae, and r2 for the current model are stored in the metrics_results dictionary, and then printed to the console, formatted to two decimal places, to give you an immediate overview of each model's performance.

MAKING PIPELINES

```
[549]: import sklearn.pipeline
from sklearn.pipeline import make_pipeline
pipe1=make_pipeline(column_transform,lr)
pipe2=make_pipeline(column_transform,dt)
pipe3=make_pipeline(column_transform,knn)
pipe4=make_pipeline(column_transform,lgb)
```

FITTING THESE PIPELINES

```
[550]: pipe1.fit(x_train,y_train)
pipe2.fit(x_train,y_train)
pipe3.fit(x_train,y_train)
pipe4.fit(x_train,y_train)
```

```
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000075 seconds.
You can set `force_row_wise=true` to remove the overhead.
And if memory is not enough, you can set `force_col_wise=true` .
[LightGBM] [Info] Total Bins 400
[LightGBM] [Info] Number of data points in the train set: 1400, number of used
features: 10
[LightGBM] [Info] Start training from score 533118.275000
/usr/local/lib/python3.12/dist-
packages/sklearn/compose/_column_transformer.py:1667: FutureWarning:
The format of the columns of the 'remainder' transformer in
ColumnTransformer.transformers_ will change in version 1.7 to match the format
of the other transformers.
At the moment the remainder columns are stored as indices (of type int). With
the same ColumnTransformer configuration, in the future they will be stored as
column names (of type str).
To use the new behavior now and suppress this warning, use
ColumnTransformer(force_int_remainder_cols=False).
```

```
warnings.warn(
```

```
[550]: Pipeline(steps=[('columntransformer',
                     ColumnTransformer(remainder='passthrough',
                                       transformers=[('onehotencoder',
```

```

OneHotEncoder(categories=[array(['Downtown', 'Rural', 'Suburban', 'Urban'],
dtype=object),
array(['No', 'Yes'], dtype=object)]),
                           ['Location', 'Garage']),
('ordinalencoder',
OrdinalEncoder(categories=[[['Excellent',
'Good',
'Fair',
'Poor']]])],
                           ['Condition']),
('standardscaler',
StandardScaler(),
['Area']))),
('lgbmregressor', LGBMRegressor()))

```

```
[551]: y_pred1=pipe1.predict(x_test)
y_pred2=pipe2.predict(x_test)
y_pred3=pipe3.predict(x_test)
y_pred4=pipe4.predict(x_test)
```

```
/usr/local/lib/python3.12/dist-packages/sklearn/utils/validation.py:2739:
UserWarning: X does not have valid feature names, but LGBMRegressor was fitted
with feature names
    warnings.warn(
```

Essentially, for each pipeline, the `x_test` data first goes through the `column_transform` (applying one-hot encoding, ordinal encoding, and scaling), and then the preprocessed data is fed into the respective regression model to produce price predictions.

CHECKING R2_SCORE

```
[552]: print(r2_score(y_test,y_pred1))
print(r2_score(y_test,y_pred2))
print(r2_score(y_test,y_pred3))
print(r2_score(y_test,y_pred4))
```

```
0.0006983615451935377
-1.146671395077441
-0.23650783832897426
-0.1739923742384868
```

```
[553]: scores=[]
for i in range(1,10000):
    x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
        ↴2,random_state=i)
    dt=DecisionTreeRegressor()
    pipe=make_pipeline(column_transform,dt)
    pipe.fit(x_train,y_train)
```

```
y_pred=pipe.predict(x_test)
scores.append(r2_score(y_test,y_pred))
```

```
[554]: np.argmax(scores)
```

```
[554]: np.int64(0)
```

```
[555]: scores[np.argmax(scores)]
```

```
[555]: -1.179504350151821
```

np.argmax(scores): This function from the NumPy library finds the index of the maximum value within the scores list. The scores list contains the R2 scores calculated during multiple iterations of training and evaluating the Decision Tree Regressor with different random train-test splits.
scores[...]: Once np.argmax(scores) returns the index of the best score, this index is then used to access the scores list itself, effectively retrieving the actual maximum R2 score that was found.

```
[556]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.
    ↪1,random_state=np.argmax(scores))
dt = DecisionTreeRegressor()
pipe=make_pipeline(column_transform,dt)
pipe.fit(x_train,y_train)
y_pred=pipe.predict(x_test)
r2_score(y_test,y_pred)
```

```
[556]: -1.0385103090629602
```

```
[557]: pipe.predict(pd.DataFrame([[1360, 3, 1970, 'Downtown', 'Excellent', 'No']], ↪
    ↪columns=['Area', 'Floors', 'YearBuilt', 'Location', 'Condition', 'Garage']))
```

```
[557]: array([149919.])
```