

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from google.colab import drive
drive.mount('/content/drive')
df=pd.read_csv('/content/drive/MyDrive/loan_data.csv')
df.head()

Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).

{"summary": "{\n    \"name\": \"df\", \n    \"rows\": 45000, \n    \"fields\": [\n        {\n            \"column\": \"person_age\", \n            \"properties\": {\n                \"dtype\": \"number\", \n                \"std\": 6.045108211348622, \n                \"min\": 20.0, \n                \"max\": 144.0, \n                \"num_unique_values\": 60, \n                \"samples\": [\n                    22.0, \n                    26.0, \n                    53.0\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\", \n                \"column\": \"person_gender\", \n                \"properties\": {\n                    \"category\", \n                    \"num_unique_values\": 2, \n                    \"samples\": [\n                        \"male\", \n                        \"female\"\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\", \n                    \"column\": \"person_education\", \n                    \"properties\": {\n                        \"category\", \n                        \"num_unique_values\": 5, \n                        \"samples\": [\n                            \"High School\", \n                            \"Doctorate\"\n                        ], \n                        \"semantic_type\": \"\", \n                        \"description\": \"\", \n                        \"column\": \"person_income\", \n                        \"properties\": {\n                            \"number\", \n                            \"std\": 80422.49863189556, \n                            \"min\": 8000.0, \n                            \"max\": 7200766.0, \n                            \"num_unique_values\": 33989, \n                            \"samples\": [\n                                48967.0, \n                                31001.0\n                            ], \n                            \"semantic_type\": \"\", \n                            \"description\": \"\", \n                            \"column\": \"person_emp_exp\", \n                            \"properties\": {\n                                \"number\", \n                                \"std\": 6, \n                                \"min\": 0, \n                                \"max\": 125, \n                                \"num_unique_values\": 63, \n                                \"samples\": [\n                                    93, \n                                    76\n                                ], \n                                \"semantic_type\": \"\", \n                                \"description\": \"\", \n                                \"column\": \"person_home_ownership\", \n                                \"properties\": {\n                                    \"category\", \n                                    \"num_unique_values\": 4, \n                                    \"samples\": [\n                                        \"OWN\", \n                                        \"OTHER\"\n                                    ], \n                                    \"semantic_type\": \"\", \n                                    \"description\": \"\", \n                                    \"column\": \"loan_amnt\", \n                                    \"properties\": {\n                                        \"number\", \n                                        \"std\": 6314.8866905411405, \n                                        \"min\": 500.0, \n                                        \"max\": 35000.0, \n                                        \"num_unique_values\": 4483, \n                                        \"samples\": [\n                                            5800.0, \n                                            28338.0\n                                        ], \n                                        \"semantic_type\": \"\", \n                                        \"description\": \"\"}\n                                }\n                            }\n                        }\n                    }\n                }\n            }\n        }\n    ]\n}
```

```

    },\n      {"column": "loan_intent",\n      "properties": {\n        "dtype": "category",\n        "num_unique_values": 6,\n        "samples": [\n          "PERSONAL",\n          "EDUCATION"],\n        "semantic_type": "\\",,\n        "description": "\\"\n      }\n    },\n    {"column": "loan_int_rate",\n    "properties": {\n      "dtype": "number",\n      "std": 2.9788082802254734,\n      "min": 5.42,\n      "max": 20.0,\n      "num_unique_values": 1302,\n      "samples": [\n        15.0,\n        13.45\n      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    },\n    {"column": "loan_percent_income",\n    "properties": {\n      "dtype": "number",\n      "std": 0.08721230801403355,\n      "min": 0.0,\n      "max": 0.66,\n      "num_unique_values": 64,\n      "samples": [\n        0.45,\n        0.54\n      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    },\n    {"column": "cb_person_cred_hist_length",\n    "properties": {\n      "dtype": "number",\n      "std": 3.8797018451620433,\n      "min": 2.0,\n      "max": 30.0,\n      "num_unique_values": 29,\n      "samples": [\n        24.0,\n        25.0\n      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    },\n    {"column": "credit_score",\n    "properties": {\n      "dtype": "number",\n      "std": 50,\n      "min": 390,\n      "max": 850,\n      "num_unique_values": 340,\n      "samples": [\n        492,\n        484\n      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    },\n    {"column": "previous_loan_defaults_on_file",\n    "properties": {\n      "dtype": "category",\n      "num_unique_values": 2,\n      "samples": [\n        "Yes",\n        "No"
      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    },\n    {"column": "loan_status",\n    "properties": {\n      "dtype": "number",\n      "std": 0,\n      "min": 0,\n      "max": 1,\n      "num_unique_values": 2,\n      "samples": [\n        0,\n        1
      ],\n      "semantic_type": "\\",,\n      "description": "\\"\n    }\n  },\n  "type": "dataframe",\n  "variable_name": "df"
}

```

## DATA PREPROCESSING AND CLEANING

```

df.shape\n(45000, 14)\n\ndf.duplicated().sum()\n\nnp.int64(0)\n\ndf.isnull().sum()

```

```

person_age          0
person_gender       0
person_education    0
person_income        0
person_emp_exp      0
person_home_ownership 0
loan_amnt          0
loan_intent         0
loan_int_rate       0
loan_percent_income 0
cb_person_cred_hist_length 0
credit_score        0
previous_loan_defaults_on_file 0
loan_status         0
dtype: int64

df.describe()

{"summary": {
    "name": "df",
    "rows": 8,
    "fields": [
        {
            "column": "person_age",
            "properties": {
                "dtype": "number",
                "std": 15895.92981900713,
                "min": 6.045108211348622,
                "max": 45000.0,
                "num_unique_values": 8,
                "samples": [
                    27.76417777777778,
                    26.0,
                    45000.0
                ],
                "semantic_type": "\",
                "description": "\n"
            }
        },
        {
            "column": "person_income",
            "properties": {
                "dtype": "number",
                "std": 2524599.5465676305,
                "min": 8000.0,
                "max": 7200766.0,
                "num_unique_values": 8,
                "samples": [
                    80319.0532222222,
                    67048.0,
                    45000.0
                ],
                "semantic_type": "\",
                "description": "\n"
            }
        },
        {
            "column": "person_emp_exp",
            "properties": {
                "dtype": "number",
                "std": 15902.409513060764,
                "min": 0.0,
                "max": 45000.0,
                "num_unique_values": 8,
                "samples": [
                    5.41033333333333,
                    4.0,
                    45000.0
                ],
                "semantic_type": "\",
                "description": "\n"
            }
        },
        {
            "column": "loan_amnt",
            "properties": {
                "dtype": "number",
                "std": 15907.41915594131,
                "min": 500.0,
                "max": 45000.0,
                "num_unique_values": 8,
                "samples": [
                    9583.157555555556,
                    8000.0,
                    45000.0
                ],
                "semantic_type": "\",
                "description": "\n"
            }
        },
        {
            "column": "loan_int_rate",
            "properties": {
                "dtype": "number",
                "std": 15906.267079441906,
                "min": 2.9788082802254734,
                "max": 45000.0,
                "num_unique_values": 8,
                "samples": [
                    11.006605777777779,
                    11.01,
                    45000.0
                ],
                "semantic_type": "\",
                "description": "\n"
            }
        }
    ]
}

```

```

\"loan_percent_income\", \n      \"properties\": {\n          \"dtype\":\n          \"number\", \n          \"std\": 15909.83858800236, \n          \"min\":\n          0.0, \n          \"max\": 45000.0, \n          \"num_unique_values\": 8, \n          \"samples\": [\n              0.1397248888888889, \n              0.12, \n              45000.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n      }, \n      {\"column\":\n      \"cb_person_cred_hist_length\", \n      \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 15907.039026136903, \n          \"min\": 2.0, \n          \"max\": 45000.0, \n          \"num_unique_values\": 8, \n          \"samples\": [\n              5.867488888888885, \n              4.0, \n              45000.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n      }, \n      {\"column\":\n      \"credit_score\", \n      \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 15718.058537708299, \n          \"min\": 50.435865000741984, \n          \"max\": 45000.0, \n          \"num_unique_values\": 8, \n          \"samples\": [\n              632.6087555555556, \n              640.0\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n      }, \n      {\"column\":\n      \"loan_status\", \n      \"properties\": {\n          \"dtype\": \"number\", \n          \"std\": 15909.819850659878, \n          \"min\":\n          0.0, \n          \"max\": 45000.0, \n          \"num_unique_values\": 5, \n          \"samples\": [\n              0.2222222222222222, \n              1.0, \n              0.41574432904844355\n          ], \n          \"semantic_type\": \"\", \n          \"description\": \"\"\n      }\n  }\n}, \n  \"type\": \"dataframe\"}

```

df.columns

```

Index(['person_age', 'person_gender', 'person_education',
'person_income',
       'person_emp_exp', 'person_home_ownership', 'loan_amnt',
'loan_intent',
       'loan_int_rate', 'loan_percent_income',
'cb_person_cred_hist_length',
       'credit_score', 'previous_loan_defaults_on_file',
'loan_status'],
      dtype='object')

```

## REMOVING UNNECESSARY COLUMNS AND CHECKING UNIQUE COLUMNS

```

df.drop(columns=['person_age', 'person_gender', 'person_education',
'person_emp_exp', 'loan_percent_income'], inplace=True)

df.rename(columns={'person_income': 'Income', 'person_home_ownership': 'O
wnership', 'loan_amnt': 'Loan Amount', 'loan_intent': 'Loan Intent
', 'loan_int_rate': 'Interest
Rate', 'cb_person_cred_hist_length': 'Credit.Hist', 'credit_score': 'Cr.Sc
ore', 'previous_loan_defaults_on_file': 'Previous
Loans', 'loan_status': 'Approval'}, inplace=True)

```

```

df.head(4)
df.describe()

{"summary": {"name": "df", "rows": 8, "fields": [
  {"column": "Income", "properties": {"dtype": "number", "std": 2524599.5465676305, "min": 8000.0, "max": 7200766.0, "num_unique_values": 8, "samples": [80319.053222222222, 67048.0, 45000.0]}, {"column": "Loan Amount", "properties": {"dtype": "number", "std": 15907.41915594131, "min": 500.0, "max": 45000.0, "num_unique_values": 8, "samples": [9583.157555555556, 8000.0, 45000.0]}, {"column": "Interest Rate", "properties": {"dtype": "number", "std": 15906.267079441906, "min": 2.9788082802254734, "max": 45000.0, "num_unique_values": 8, "samples": [11.006605777777779, 11.01, 45000.0]}, {"column": "Credit.Hist", "properties": {"dtype": "number", "std": 15907.039026136903, "min": 2.0, "max": 45000.0, "num_unique_values": 8, "samples": [5.8674888888888885, 4.0, 45000.0]}, {"column": "Cr.Score", "properties": {"dtype": "number", "std": 15718.058537708299, "min": 50.435865000741984, "max": 45000.0, "num_unique_values": 8, "samples": [632.608755555556, 640.0, 45000.0]}, {"column": "Approval", "properties": {"dtype": "number", "std": 15909.819850659878, "min": 0.0, "max": 45000.0, "num_unique_values": 5, "samples": [0.2222222222222222, 1.0, 0.41574432904844355]}]}, "type": "dataframe"}}

df['Ownership'].unique()
array(['RENT', 'OWN', 'MORTGAGE', 'OTHER'], dtype=object)

df['Approval'].unique()
array([1, 0])

```

```
df['Previous Loans'].unique()

array(['No', 'Yes'], dtype=object)

df['Cr.Score'].unique()

array([561, 504, 635, 675, 586, 532, 701, 585, 544, 640, 621, 651,
573,
    708, 583, 670, 663, 694, 709, 679, 684, 662, 691, 600, 654,
626,
    607, 700, 553, 589, 681, 567, 669, 606, 582, 649, 602, 616,
631,
    637, 695, 620, 622, 645, 624, 570, 648, 652, 559, 623, 609,
579,
    688, 661, 562, 664, 564, 598, 557, 677, 690, 599, 604, 601,
634,
    671, 789, 538, 587, 683, 518, 617, 668, 673, 706, 536, 689,
595,
    584, 642, 614, 597, 625, 603, 643, 508, 505, 593, 686, 646,
697,
    615, 687, 650, 588, 658, 531, 665, 703, 594, 618, 574, 577,
653,
    630, 660, 639, 612, 628, 592, 580, 678, 672, 613, 566, 718,
484,
    699, 656, 659, 636, 554, 807, 578, 674, 608, 569, 629, 560,
548,
    667, 676, 581, 655, 551, 529, 666, 576, 633, 611, 657, 647,
542,
    692, 545, 540, 525, 537, 641, 539, 563, 712, 491, 590, 572,
528,
    638, 627, 596, 547, 507, 565, 693, 522, 632, 556, 499, 704,
503,
    714, 552, 555, 558, 521, 605, 571, 591, 719, 610, 535, 644,
523,
    546, 702, 711, 534, 805, 682, 447, 725, 680, 568, 524, 710,
707,
    619, 460, 696, 527, 511, 477, 575, 496, 685, 476, 502, 541,
722,
    506, 487, 530, 515, 513, 520, 724, 514, 501, 549, 720, 486,
716,
    465, 509, 517, 444, 482, 698, 443, 512, 526, 454, 550, 462,
717,
    727, 715, 543, 448, 713, 721, 481, 723, 533, 705, 483, 740,
495,
    459, 489, 519, 494, 498, 488, 468, 490, 479, 510, 472, 480,
421,
    473, 500, 493, 441, 485, 450, 467, 461, 475, 463, 728, 735,
458,
    456, 726, 729, 750, 455, 478, 453, 445, 431, 737, 497, 492,
437,
```


```
    516, 451, 434, 435, 439, 466, 440, 449, 469, 471, 730, 470,
734,
    751, 739, 736, 738, 733, 732, 731, 741, 748, 744, 762, 742,
745,
    743, 464, 390, 755, 747, 759, 746, 756, 753, 764, 850, 760,
754,
    770, 784, 773, 765, 772, 419, 474, 430, 418, 420, 792, 768,
457,
    446, 767])
df['Credit.Hist'].unique()

array([ 3.,  2.,  4.,  8.,  7.,  6.,  9., 10.,  5., 11., 16., 15.,
12.,
   13., 17., 14., 25., 28., 27., 22., 19., 29., 23., 26., 20.,
21.,
   30., 24., 18.])
```

CHECKING CO-RELATION% BETWEEN COLUMNS

```
plt.figure(figsize=(8,6))
plt.title('Heatmap for co-relation')
sns.heatmap(df.corr(numeric_only=True), annot=True)

<Axes: title={'center': 'Heatmap for co-relation'}>
```




Loan amount has higher degree of relation with income, and approval of loan has high degree relation with interest rate

#### TRAIN\_TEST\_SPLIT\_

```
df.columns
Index(['Income', 'Ownership', 'Loan Amount', 'Loan Intent ', 'Interest
Rate',
       'Credit.Hist', 'Cr.Score', 'Previous Loans', 'Approval'],
      dtype='object')

sns.boxplot(df['Income'])

<Axes: ylabel='Income'>
```



```
df['Income'].value_counts()

Income
8000.0    15
73011.0   10
36995.0   9
37020.0   8
60914.0   8
...
16720.0    1
29547.0    1
64117.0    1
89930.0    1
42960.0    1
Name: count, Length: 33989, dtype: int64

len(df[df['Income']>400000])

165


df=df[df['Income']<400000]

df.shape

(44835, 9)

sns.boxplot(df['Loan Amount'])
```

```
<Axes: ylabel='Loan Amount'>
```




```
df['Loan Amount'].value_counts()
```

```
Loan Amount
10000.0    3607
5000.0     2786
6000.0     2423
12000.0    2412
15000.0    1999
...
6074.0      1
3108.0      1
2457.0      1
4601.0      1
17783.0     1
Name: count, Length: 4451, dtype: int64
```

```
sns.boxplot(df['Interest Rate'])
```

```
<Axes: ylabel='Interest Rate'>
```



```
df['Interest Rate'].value_counts()
```

Interest Rate	Count
11.01	3318
10.99	802
7.51	796
7.49	686
7.88	671
...	
19.25	1
19.62	1
19.11	1
19.80	1
16.92	1

Name: count, Length: 1302, dtype: int64

```
x=df.drop(columns=['Approval'])
x
```

```
{"summary": {"name": "x", "rows": 44835, "fields": [
  {"column": "Income", "properties": {"dtype": "number", "std": 45588.23413018292, "min": 8000.0, "max": 397048.0, "num_unique_values": 33824, "samples": [27049.0, 24652.0, 79904.0]}, "semantic_type": "\\", "description": "\"\\n\""}, {"column": "Ownership", "description": "\"\\n\""}]}
```

```
\"properties\": {\n    \"dtype\": \"category\", \n    \"num_unique_values\": 4, \n    \"samples\": [\n        \"OWN\", \n        \"OTHER\", \n        \"RENT\"\n    ], \n}, \n    \"semantic_type\": \"\", \n    \"column\": \"Loan Amount\", \n}, \n    \"properties\": {\n        \"dtype\": \"number\", \n        \"min\": 500.0, \n        \"max\": 6290.258971072717, \n        \"num_unique_values\": 4451, \n        \"samples\": [\n            17507.0, \n            14807.0, \n            6016.0\n        ], \n        \"semantic_type\": \"\", \n        \"description\": \"\\n        }\\n    }, \n        {\\n            \"column\": \"Loan Intent\", \n            \"properties\": {\n                \"category\": \"\", \n                \"num_unique_values\": 6, \n                \"samples\": [\n                    \"PERSONAL\", \n                    \"EDUCATION\", \n                    \"DEBTCONSOLIDATION\"\n                ], \n                \"semantic_type\": \"\", \n                \"description\": \"\\n                }\\n            }, \n            {\\n                \"column\": \"Interest Rate\", \n                \"properties\": {\n                    \"number\": \"\", \n                    \"std\": 2.9794284164692604, \n                    \"min\": 5.42, \n                    \"max\": 20.0, \n                    \"num_unique_values\": 1302, \n                    \"samples\": [\n                        7.8, \n                        8.12, \n                        16.96\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\\n                }\\n            }, \n            {\\n                \"column\": \"Credit.Hist\", \n                \"properties\": {\n                    \"dtype\": \"number\", \n                    \"std\": 3.859562852100272, \n                    \"min\": 2.0, \n                    \"max\": 30.0, \n                    \"num_unique_values\": 29, \n                    \"samples\": [\n                        24.0, \n                        27.0, \n                        12.0\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\\n                }\\n            }, \n            {\\n                \"column\": \"Cr.Score\", \n                \"properties\": {\n                    \"number\": \"\", \n                    \"std\": 50, \n                    \"min\": 390, \n                    \"max\": 807, \n                    \"num_unique_values\": 337, \n                    \"samples\": [\n                        652, \n                        522, \n                        557\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\\n                }\\n            }, \n            {\\n                \"column\": \"Previous Loans\", \n                \"properties\": {\n                    \"category\": \"\", \n                    \"num_unique_values\": 2, \n                    \"samples\": [\n                        \"Yes\", \n                        \"No\"\n                    ], \n                    \"semantic_type\": \"\", \n                    \"description\": \"\\n                }\\n            }\n        },\n        \"type\": \"dataframe\", \n        \"variable_name\": \"x\"\n    }\n}
```

```
y=df['Approval']
```

y

```

44997    1
44998    1
44999    1
Name: Approval, Length: 44835, dtype: int64

from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3,random_state=42)

```

## ENCODING DATA

```

from sklearn.preprocessing import OneHotEncoder
ohe=OneHotEncoder()
ohe.fit(x[['Ownership','Loan Intent']])

OneHotEncoder()

from sklearn.preprocessing import LabelEncoder
le=LabelEncoder()
x['Previous Loans']=le.fit_transform(x['Previous Loans'])

```

## DATA SCALING

```

from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(x[['Interest Rate']])
scaler.fit(x[['Income']])
scaler.fit(x[['Loan Amount']])

StandardScaler()

from math import remainder
from sklearn.compose import make_column_transformer
from sklearn.pipeline import make_pipeline
column_transform=make_column_transformer((OneHotEncoder(categories=ohe.categories_),['Income','Loan Intent']),
                                         (StandardScaler(),['Interest Rate','Loan Amount','Income']),
                                         remainder='passthrough')
column_transform

ColumnTransformer(remainder='passthrough',
                  transformers=[('onehotencoder',
                                 OneHotEncoder(categories=[array(['MORTGAGE', 'OTHER', 'OWN', 'RENT'],
                                 dtype=object),
                                 array(['DEBTCONSOLIDATION', 'EDUCATION', 'HOMEIMPROVEMENT', 'MEDICAL',
                                 'PERSONAL', 'VENTURE'], dtype=object))],
                  ['Income', 'Loan Intent']),

```

```

        ('standardscaler', StandardScaler(),
         ['Interest Rate', 'Loan Amount',
          'Income'])))

from sklearn.preprocessing import OneHotEncoder, StandardScaler,
OrdinalEncoder
from sklearn.compose import make_column_transformer
from sklearn.pipeline import make_pipeline

# Define categorical columns for OneHotEncoder, numerical columns for
StandardScaler,
# and the binary categorical column for OrdinalEncoder.
categorical_cols_ohe = ['Ownership', 'Loan Intent']
numerical_cols = ['Income', 'Loan Amount', 'Interest Rate',
'Credit.Hist', 'Cr.Score']
binary_categorical_col = ['Previous Loans'] # This column needs
explicit encoding

# Create the ColumnTransformer
column_transform = make_column_transformer(
    (OneHotEncoder(categories=ohe.categories_), categorical_cols_ohe),
    (StandardScaler(), numerical_cols),
    (OrdinalEncoder(), binary_categorical_col), # OrdinalEncoder will
convert 'Yes'/ 'No' to numerical
    remainder='drop' # Explicitly drop any unhandled columns
)

x_train_trans=column_transform.fit_transform(x_train)
x_test_trans=column_transform.transform(x_test)

```

## MODELING

```

#LINEAR REGRESSION (changed to Logistic Regression for classification)
from sklearn.linear_model import LogisticRegression
lr=LogisticRegression(max_iter=1000) # Increased max_iter for
convergence

#TREE BASED REGRESSION (changed to Decision Tree Classifier)
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier()

#DISTANCE BASED MODELS (changed to K-Nearest Neighbors Classifier)
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier()

#RANDOM FOREST CLASSIFIER
from sklearn.ensemble import RandomForestClassifier # Changed from
Regressor to Classifier
rfc=RandomForestClassifier()

```

```

#BAGGING CLASSIFIER
from sklearn.ensemble import BaggingClassifier
bgc=BaggingClassifier()

#EXTRA TREE CLASSIFIER
from sklearn.ensemble import ExtraTreesClassifier
etc=ExtraTreesClassifier()

#XGB CLASSIFIER
from xgboost import XGBClassifier
xgb=XGBClassifier()

#METRICS
from sklearn.metrics import
accuracy_score,precision_score,recall_score,f1_score,confusion_matrix,
ConfusionMatrixDisplay

```

Linear Regression: It imports LinearRegression for linear modeling. Tree-Based Regression: It imports DecisionTreeRegressor for decision tree-based regression tasks. Distance-Based Models: KNeighborsRegressor is imported, which is a distance-based algorithm for regression. Ensemble Models: It imports several ensemble methods: RandomForestRegressor, BaggingClassifier, and ExtraTreesClassifier. These models combine predictions from multiple base estimators to improve robustness and accuracy. XGBoost Classifier: XGBClassifier is imported from the xgboost library, which is a powerful gradient boosting framework for classification. Metrics: Finally, it imports a comprehensive set of evaluation metrics from sklearn.metrics, including accuracy\_score, precision\_score, recall\_score, f1\_score, confusion\_matrix, and ConfusionMatrixDisplay. These will be used to assess the performance of the trained models.

```

classifiers={
    'LR':lr,
    'DT':dt,
    'KNN':knn,
    'RFC':rfc,
    'XGB':xgb
}

def train_classifier(classifier,x_train,x_test,y_train,y_test):

    #TRAINING CLASSIFIER
    classifier.fit(x_train,y_train)

    #MAKING PREDICTIONS
    y_pred=classifier.predict(x_test)

    #CALCULATING METRICS
    accuracy=accuracy_score(y_test,y_pred)
    precision=precision_score(y_test,y_pred,average='weighted')
    recall=recall_score(y_test,y_pred,average='weighted')

```

```

f1=f1_score(y_test , y_pred, average='weighted')

#GENERATING AND DISPLAYING CONFUSION MATRIX
cm=confusion_matrix(y_test,y_pred)
disp=ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(cmap=plt.cm.Blues)

plt.title(f'Confusion Matrix for {name}')
plt.xlabel("PREDICTED LABELS")
plt.ylabel("TRUE LABELS")
plt.show()

return accuracy,precision,recall,f1

```

The code in this cell defines a Python function called `train_classifier` which automates the process of training a classification model, making predictions, evaluating its performance, and visualizing the results. Let's break it down:

**Function Definition:** The function `train_classifier` takes five arguments: `classifier` (the machine learning model you want to train), `x_train` and `y_train` (your training data and labels), and `x_test` and `y_test` (your testing data and true labels). **Training Classifier:** `classifier.fit(x_train, y_train)` trains the provided classifier model using the training features (`x_train`) and their corresponding true labels (`y_train`). The model learns patterns from this data. **Making Predictions:** `y_pred=classifier.predict(x_test)` uses the now-trained model to make predictions on the unseen test features (`x_test`). The results are stored in `y_pred`. **Calculating Metrics:** It then calculates several key classification metrics: `accuracy_score`: The proportion of correctly classified instances. `precision_score`: The ability of the classifier not to label as positive a sample that is negative (weighted average for multi-class). `recall_score`: The ability of the classifier to find all the positive samples (weighted average for multi-class). `f1_score`: The harmonic mean of precision and recall (weighted average for multi-class). **Generating and Displaying Confusion Matrix:** A `confusion_matrix` is created to summarize the performance of a classification algorithm. `ConfusionMatrixDisplay` then visualizes this matrix, showing the counts of true positive, true negative, false positive, and false negative predictions. **Plot Customization:** It sets a title for the confusion matrix plot, labels the x and y axes as 'PREDICTED LABELS' and 'TRUE LABELS' respectively, and finally `plt.show()` displays the plot. **Return Values:** The function returns the calculated accuracy, precision, recall, and f1 scores.

```

# STORING ACCURACY, PRECISION, RECALL, AND F1 SCORES
accuracy_scores=[]
precision_scores=[]
recall_scores=[]
f1_scores=[]

#LOOP THROUGH EACH CLASSIFIER
for name,clf in classifiers.items():
    print(f'\nTraining {name}...')
    current_accuracy, current_precision, current_recall, current_f1 =
train_classifier(clf, x_train_trans, x_test_trans, y_train, y_test)

```


```

# PRINT THE CLASSIFIER'S PERFORMANCE MATRIX
print(f'Name: {name}')
print(f'Accuracy: {current_accuracy}')
print(f'Precision: {current_precision}')
print(f'Recall: {current_recall}')
print(f'F1 Score: {current_f1}')
print('-----')

# APPEND SCORES TO LIST
accuracy_scores.append(current_accuracy)
precision_scores.append(current_precision)
recall_scores.append(current_recall)
f1_scores.append(current_f1)

```

Training LR...




```

Name: LR
Accuracy: 0.8922756672366366
Precision: 0.8923673362820014
Recall: 0.8922756672366366
F1 Score: 0.8923211892149823

```


-----  
Training DT...



Name: DT  
Accuracy: 0.9020147200951603  
Precision: 0.9029994148063958  
Recall: 0.9020147200951603  
F1 Score: 0.9024678804882862  
-----

Training KNN...

Confusion Matrix for KNN




Name: KNN  
Accuracy: 0.9002304661363467  
Precision: 0.8977673053331857  
Recall: 0.9002304661363467  
F1 Score: 0.8985125067472145

-----

Training RFC...


Confusion Matrix for RFC



Name: RFC  
Accuracy: 0.9295219686268679  
Precision: 0.9282512348781771  
Recall: 0.9295219686268679  
F1 Score: 0.9285395966858785

-----

Training XGB...



```
Name: XGB
Accuracy: 0.9332391643743959
Precision: 0.932050963629974
Recall: 0.9332391643743959
F1 Score: 0.9322543471153916
```

First, it sets up four empty lists: `accuracy_scores`, `precision_scores`, `recall_scores`, and `f1_scores`. These lists will store the performance metrics for each classifier after it's been trained and evaluated.

**Looping Through Classifiers:** The `for name, clf in classifiers.items():` loop iterates through each classifier in your `classifiers` dictionary. In each iteration:

`name` will be the string key (e.g., 'LR', 'DT', 'XGB'). `clf` will be the actual classifier object (e.g., `LogisticRegression()`, `DecisionTreeClassifier()`). **Training and Evaluation:** Inside the loop, it calls the `train_classifier` function (which was defined in the previous cell). This function:

Trains the current `clf` using `x_train_trans` and `y_train`. Makes predictions on `x_test_trans`. Calculates the accuracy, precision, recall, and `f1_score` using `y_test` and the predictions. Displays a confusion matrix for the current classifier. Returns these four performance metrics. **Printing Results:** After the `train_classifier` function returns the scores, the code prints the Name of the

classifier and its calculated Accuracy, Precision, Recall, and F1 Score to the console, providing a clear overview of how each model performed.

Storing Scores: Finally, accuracy\_scores.append(current\_accuracy) (and similarly for precision, recall, and f1) adds the scores of the current classifier to their respective lists. This accumulates the results, allowing you to later compare the performance of all classifiers collectively.

```
results = {name: acc for name, acc in zip(classifiers.keys(),  
accuracy_scores)}  
best_model_name = max(results, key=results.get)  
print("Best Model:", best_model_name)  
  
Best Model: XGB
```

## TRAINING PIPELINE

```
from sklearn.pipeline import Pipeline  
  
final_model = Pipeline(steps=[  
    ("preprocessing", column_transform),  
    ("model", classifiers[best_model_name])  
])  
final_model.fit(x_train,y_train)  
  
Pipeline(steps=[('preprocessing',  
    ColumnTransformer(transformers=[('onehotencoder',  
        OneHotEncoder(categories=[array(['MORTGAGE', 'OTHER', 'OWN', 'RENT'],  
        dtype=object),  
  
        array(['DEBTCONSOLIDATION', 'EDUCATION', 'HOMEIMPROVEMENT', 'MEDICAL',  
        'PERSONAL', 'VENTURE'], dtype=object)),  
        ['Ownership',  
         'Loan Intent ']),  
        ('standardscaler',  
        StandardScaler(),  
        ['Income', 'Loan  
Amoun...  
        feature_types=None,  
        feature_weights=None,  
        gamma=None, grow_policy=None,  
        importance_type=None,  
        interaction_constraints=None,  
        learning_rate=None,  
        max_bin=None, max_cat_threshold=None,  
        max_cat_to_onehot=None,  
        max_delta_step=None,  
        max_depth=None, max_leaves=None,  
        min_child_weight=None, missing=nan,  
        monotone_constraints=None,
```

```
multi_strategy=None,  
n_estimators=None, n_jobs=None,  
num_parallel_tree=None, ...))))
```

## REAL LIFE PREDICTION

```
new_person = pd.DataFrame([  
    {  
        "Income": 85000,  
        "Ownership": "RENT",  
        "Loan Amount": 15000,  
        "Loan Intent": "PERSONAL",  
        "Interest Rate": 12.5,  
        "Credit.Hist": 4,  
        "Cr.Score": 690,  
        "Previous Loans": "No"  
    }  
)  
  
prediction = final_model.predict(new_person)  
if prediction[0] == 1:  
    print("Loan Approved")  
else:  
    print("Loan Rejected")  
  
Loan Rejected
```